

FROM PIE IN THE SKY TO ROI

GETTING A.I. DEPLOYMENT RIGHT IN THE
FINANCIAL SERVICES INDUSTRY

JULY 2026

QUINLAN
& ASSOCIATES

TABLE OF CONTENTS

SECTION	TITLE	PAGE
SECTION 1	HIGH INVESTMENT, LOW RETURN	4
SECTION 2	STRATEGISE	13
SECTION 3	GOVERN	19
SECTION 4	ENABLE	23
SECTION 5	HOW WE CAN HELP	41

EXECUTIVE SUMMARY

The financial services industry is entering a new phase of artificial intelligence (“A.I.”) adoption, with global financial institutions (“FIs”) investing an estimated USD 76 billion in A.I.-related initiatives in 2025, a number we forecast to reach USD 132 billion by 2030.


Despite vast investments and institutional ambition, most financial institutions struggle to translate A.I. initiatives into measurable business value and clear returns on investment. In many cases, deployment activity has outpaced the development of the organisational foundations required to support it. Strategy is frequently backward engineered around available technology, risk is considered only after design and development have progressed, and existing operating models are not equipped to govern, implement, and scale A.I. effectively.

To overcome these barriers, financial institutions should adopt an integrated roadmap across three interconnected layers:

- **Strategise:** A.I. should be treated as an enabling capability to deliver the enterprise strategy rather than as a standalone technology agenda. Institutions should balance outside-in awareness of emerging technological possibilities with an inside-out assessment of strategic priorities, customer journeys, and operational workflows. This requires systematically identifying friction points across existing processes, determining where A.I. can generate meaningful value, and applying non-negotiable eligibility thresholds before prioritising viable use cases. Beyond improvements to existing processes, there is a broader opportunity for institutions to fundamentally reimagine workflows around A.I. capabilities, which can potentially unlock new ways of delivering value.
- **Govern:** Governance should be established before material resources are committed. Traditional approaches that treat governance as a downstream approval process are proving insufficient for dynamic A.I. environments. Financial institutions should implement systematic eligibility gates to exclude use cases that do not meet minimum viable requirements. Governance must also extend beyond use-case approval into the underlying workflow, defining how A.I. operates within controlled boundaries through runtime controls, human checkpoints, and more. When designed effectively, early governance can serve not only as a compliance safeguard, but also as a competitive edge for more effective and confident deployment by establishing clear boundaries.
- **Enable:** Once strategic priorities and permissible guardrails have been established, institutions must develop the technological and operational capabilities required for execution. From a technology perspective, this includes selecting an appropriate model architecture and developing the surrounding A.I. engineering capabilities, including prompt engineering, context management, and harnesses to deliver reliable, secure, and cost-effective outcomes. Enablement also requires a fundamental shift in people, roles, and ways of working. As A.I. becomes embedded across institutional workflows, frontline staff will increasingly direct A.I. outputs, middle managers will orchestrate mixed teams of humans and A.I. agents, and senior leaders will be responsible for setting strategic direction, boundaries, and organisational design. A.I. transformation therefore cannot rely solely on a traditional top-down technology rollout as meaningful value is often identified closer to the frontline, where real domain knowledge and operational realities are most visible.

Ultimately, the quality of transformation will depend on how effectively institutions codify human expertise into the context, rules, controls, and validation logic that guide A.I.-enabled workflows. Technology can automate and augment execution, but it cannot independently determine how an institution should operate, which risks it should accept, or what constitutes an appropriate outcome. Those decisions must remain grounded in institutional strategy, domain expertise, and clearly assigned accountability.

As financial institutions continue to advance their A.I. ambitions, particularly towards more autonomous and agentic models, they should conduct a structured readiness diagnostic covering strategy, governance, technology, people, and processes. This diagnostic can identify critical gaps, prioritise interventions, and establish a practical 12 to 24-month roadmap for scalable adoption. By strengthening these foundations before accelerating deployment, institutions can shift from fragmented experimentation towards controlled, repeatable, and measurable value creation.



SECTION 1

LARGE INVESTMENTS, DISAPPOINTING RETURNS

QUINLAN
& ASSOCIATES

GLOBAL A.I. SPENDING ACROSS THE FINANCIAL SERVICES INDUSTRY


The global financial services industry invested USD 75 billion in A.I.-related initiatives in 2025, with this figure projected to reach USD 132 billion by 2030; APAC commitment and allocation to A.I. continue to lag their peers in North America and Europe

Global A.I. Spending Size

Financial Services, USD Billion, 2025-30E


1.2 – 2.7x

Expected increase in A.I.-related budget allocation over the next five years



Average A.I. Spending Allocation

Financial Services, % A.I. of IT Expenditure, 2025



¹Retail, commercial, and corporate banks; ²Investment banks and brokerages; ³Asset managers, wealth managers, private banks, and high frequency trading firms; ⁴Securities exchanges, digital assets exchanges, and payment providers; ⁵New initiatives, including AI




Source: Public disclosures (e.g., annual statements, financial positions, etc.), Expert interviews, Oxford Economics, Coalition Greenwich, Nvidia, Quinlan & Associates estimation

A.I. DEPLOYMENT AREAS

Across banking, capital markets, and insurance, A.I. is increasingly viewed as a source of business value across both external client-facing propositions and internal operating functions

Current and Expected Deployment Areas

By Sector, 2025




¹Average of Personal and Commercial Lines

STATE OF A.I. ADOPTION IN FINANCIAL SERVICES

While financial institutions are accelerating A.I. deployment, adoption remains concentrated in assistive use cases, including chatbots and copilots, while agentic A.I. has yet to achieve widespread adoption

Types of A.I. Applications

By Level of Complexity & Adoption




¹Asset managers; ²Examples are non-exhaustive

WIDELY ADOPTED YET CHALLENGING TO CAPTURE VALUE

More than 70% of traditional financial institutions struggle to translate A.I. investment into measurable value and clear ROI, often because strategy is backward engineered around technology, risk is considered too late, and the operating model is not fit for purpose

Adoption Status of A.I.

2026, n=149, Traditional financial institutions across all sizes




Most financial institutions are actively adopting A.I. Solutions

VS.

Challenging to Capture A.I. Value

2026, n=101, Traditional financial institutions across all sizes



Yet more than 2/3 are having difficulties in capturing the A.I. Value

- 1 Strategy is backward engineered**
 Organisations often select A.I. tools before defining the business problems, workflows, and outcomes they should address. Strategy is then adapted to fit the technology, rather than technology being selected to support business priorities
- 2 Risk is treated as an afterthought**
 Risk is often assessed after use cases and solutions have been selected, rather than being embedded into their design. This creates late-stage remediation, inconsistent controls, and unclear ownership across the A.I. lifecycle
- 3 Operating models are not fit-for-purpose**
 Existing operating models often lack the proper governance, measurement frameworks, workforce capabilities, and change-management disciplines required to scale A.I. As a result, A.I. fails to be embedded across the organisation

1. STRATEGY IS BACKWARD ENGINEERED

The rush to deploy A.I. in response to ‘fear of missing out’ can generate visible momentum and media attention, but initiatives launched without strategy-led design often fail to improve underlying business performance

Problem

Backward Engineering



2. RISK IS TREATED AS AN AFTERTHOUGHT

As LLM-based solutions ingest increasingly sensitive data and assume greater autonomy (e.g., recommending actions or making decisions on users' behalf), risk must be embedded early to avoid control gaps, costly remediation, and project cancellation

Technology Deployment Journey

Problem

IDEATION

Map out business requirements, which are used to determine the core functional requirements of the prospective solution

Proposition Development

Determining the specific value-add of the suggested solution

Features Identification

Laying out features of the solution that address the specific business problem

UI / UX Mapping & Verification

Conceptualising the user experience and aligning it with end-users

VALIDATION

Demonstrate that the solution concepts have the potential to be applied and carried out in a real-world situation

Feasibility Assessment

Understanding capability gaps that are required to achieve the future state

Resource Requirement

Proposing timeline, budget, and human resources required for the initiative

Risk Identification

Identifying potential risks that may arise prior to development and implementation

DEVELOPMENT

Carry out the development of functionalities and test it out with the end users to collect and incorporate feedback, iteratively

Prototype Development

Building out core functionalities that aim to address the business needs

Agile Feedback Loop


Seeking and incorporating the end-users' feedback under a controlled environment

System Integration

Integrating with existing systems and deploying the solution

The Reality of A.I. Adoption

% of A.I. Initiatives¹, Dec 2025



¹Out of 25,289 A.I. initiatives developed by 400 surveyed banks; ²1,649 business leaders worldwide across 14 sectors, including 15% in banking and 4% in insurance

3. OPERATING MODELS ARE NOT FIT FOR PURPOSE

Across the A.I. lifecycle, different blockers emerge at different stages, but they consistently trace back to one structural root cause: the operating model is no longer equipped to support A.I. adoption at scale

A.I. Lifecycle / Value Chain

Operating Model-related Issues, Illustrative

	DATA INGESTION	WORKFLOW DEFINITION	MODEL SELECTION	SOLUTION DEPLOYMENT	OUTPUT EVALUATION	SCALING OPERATION
Common Problems Among Institutions	Fragmented and Unreliable Inputs A.I. systems may ingest inconsistent, outdated, or over-permissioned data across repositories, resulting in hallucination risks	Persistent Workflow Inefficiencies Organisations fail to unlock the full potential of A.I., as workflows are inefficiently adapted around A.I. capabilities (vs. redesigned to address business needs)	Trade-off in Capability vs. Explainability Default to less capable, pre-approved models, prohibiting the use of outperforming, frontier models at the outset due to model risk requirements	Fragmented A.I. Deployment When capability is handed to different functions, teams assemble their own agents, resulting in duplicative efforts and fragmented deployment	Inconsistent A.I. Output Quality A.I. outputs are generated efficiently but lack business value due to inconsistent evaluation processes and unclear quality benchmarks	Limited Enterprise-Wide A.I. Adoption A.I. tools remain underutilised across the organisation, preventing widespread adoption and limiting the return on investment
Current State of Operating Model	Poor Data Governance - Customer data sitting in multiple repositories across business lines (e.g., private banking and wealth) - Inconsistent ownership and classification of data	Technology-first Approach A.I. is deployed as a solution looking for a problem, resulting in workflows being reshaped around A.I. vs. being purpose-built around workflow requirements	Lack of Robust Evaluation - Model selection is driven primarily by risk avoidance and binary approval mechanisms - The same model is evaluated and applied the same way across use cases	Decentralised Build - A lack of central coordination where development and deployment of A.I. solutions across functions are done in isolation - Limited interoperability	No Shared Standards Lack of standardised workflows, human review checkpoints, and use-case specific KPIs to assess and validate A.I. outputs	Limited Capability Building - Existing training and change management efforts have not evolved alongside changing employee roles - Staff do not have the skills and confidence to use the A.I. tools

Current operating models are no longer fit to support A.I. adoption at scale – addressing these problems is therefore important for:

- Faster Issue Resolution
- Reduced Operational, Regulatory, and Model-related Risks¹
- Improved Compliance
- Accelerated Time to Market


¹Includes hallucination risks, data leakage, etc.

BUILDING BLOCKS FOR A.I. DEPLOYMENT

To overcome these barriers and realise value at scale, financial institutions must address three interconnected layers that collectively form a practical roadmap for A.I. adoption: strategise, govern, and enable

Building Blocks

Strategy & Operating Model



STRATEGISE

Define Objective

Is our A.I. strategy and plan surrounding the use cases and models applied to the right problems, and do we have a clear line of sight from use cases to business outcomes?

DESCRIPTION

- Translate overarching business objectives into a prioritised portfolio of use cases
- Identify use cases through workflow-led assessment
- Adopt a structured prioritisation framework rather than relying on ad-hoc selection

GOVERN

Establish Controls

Is A.I. pre-selected and operating within defined risk parameters and delivering reliable, high-quality outputs by following appropriate controls?

- Establish gating criteria (e.g., data readiness and risk thresholds) prior to use case prioritisation
- Centralise governance backed by a systematic playbook with clear accountability
- Redesign controls for A.I.-enabled workflows and technology operations, executed by teams
- Embed KPIs across the end-to-end A.I. lifecycle

ENABLE

Drive Better Usage

Do we have the appropriate capabilities in place to adopt A.I. at scale and in a way that sustains value over time?

- Select the right models and approaches for each use case, not just the most capable
- Build A.I. literacy across the organisation, from leaders to front-line staff, to reflect their evolving roles
- Establish ways of working that embed A.I. and governance into day-to-day processes, not as an add-on



SECTION 2

STRATEGISE


QUINLAN
& ASSOCIATES

DEFINING THE A.I. STRATEGY

A.I. value is maximised when organisations balance outside-in awareness of emerging possibilities with a rigorous inside-out approach that embeds A.I. into core workflows and strategy-led priorities

A.I. Strategy

Double-edge Perspectives



This ensures that market-proven A.I. capabilities are selectively embedded into workflows and operating models to generate actual value

STRATEGY-FIRST APPROACH

Financial institutions should first determine whether and where A.I. supports the overarching business strategy, establishing clear directional intent before identifying use cases where adoption can create meaningful value

Strategic Alignment

Existing Strategy and A.I. Capabilities (Illustrative)

Type of Institutions	 RETAIL / COMMERCIAL BANKS	 ASSET / WEALTH MANAGERS	 SECURITIES BROKERAGES	 INSURANCE FIRMS
Strategic Objective (Illustrative)	Rationalise the Footprint Streamline the bank's structure and operations by exiting non-core markets and service lines	Position for Wealth Transfer Ensure smooth asset transition and reinvent current workflows to align with the values and behaviours of the next generation	Deepen Retail Penetration Deepen monetisation avenues with retail clients by reinventing the "do-it-yourself" approach to investment	Expand Protection Services Introduce innovative products to protect clients and their assets from a range of emerging risks

	IDENTIFY EMERGING RISKS	DEVELOP BUSINESS CASES	DESIGN PRODUCT & PRICING	OFFLOAD AGGREGATE RISKS	LAUNCH & MARKET PRODUCTS	PROCESS FUTURE CLAIMS
Journey	Pinpoint market gaps and test if the new risks are insurable	Build formal business cases to establish the target market	Set premium pricing and write the policy guidelines	Manage the insurer's exposure and obtain regulatory approvals	Train agents and brokers for marketing and process claims	Set up robust workflows to review and pay out claims
Relevant A.I. Capabilities (Example)	Horizon Scanning A.I. can continuously scrape unstructured data (e.g., litigation, research) for signals of new risks	Persona Testing A.I. can simulate client personas and stress-test if and how the product may be exploited	Policy Drafting A.I. can draft precise definitions and exclusions before comparing them with existing policies	Pitch Creation A.I. can create collateral, such as impact projects, that are presented to reinsurers and regulators	Sales Enablement A.I. chatbots and employee copilots can inform customers and agents more intuitively	Pre-processing A.I. can automate document processing and provide preliminary opinions on each claim
Alignment Assessment	Highly Aligned	Somewhat Aligned	Highly Aligned	Somewhat Aligned	Highly Aligned	Highly Aligned

WORKFLOW DESIGN (1/2) – REVIEW EXISTING PROCESSES

Within the relevant functional areas, institutions should identify operational frictions and bottlenecks across end-to-end workflows, establishing a structured basis for applying A.I. and redirecting resources towards higher-value activities

Use Case Development

A.I. Deployment

	Description	Implications of Inaction
DEFINE	STANDARD OPERATING PROCEDURE / WORKFLOW Map the end-to-end workflow across business lines, including inputs, outputs, decision points, approvals, controls, and hand-offs	<i>Automating fragmented, undocumented processes, resulting in inefficiencies and weakened controls</i>
ASSESS	MANUAL GRIND AND PAIN POINTS Identify workflow bottlenecks and manual-intensive activities by assessing frequency, effort, turnaround time, rework, error rates, and control sensitivity	<i>Deploying A.I. based on assumptions or hype rather than real pain points and business needs that generate value</i>
MATCH	RISK SUITABILITY ASSESSMENT Evaluate whether each activity is suitable for automation, augmentation, or continued human execution, taking into account regulatory requirements and model risk / impact	<i>Applying A.I. to unsuitable use cases where risks exceed potential value / benefits (e.g., compliance, op. risks, etc.)</i>
REDESIGN	A.I.-INTEGRATED WORKFLOW Revisit and update the end-to-end workflow to incorporate A.I., defining new roles, controls, escalation paths, and oversight as needed	<i>Resources intended to be 'freed up' end up correcting A.I.'s mistakes due to an unclear, outdated process</i>

Manual vs. High-Value Tasks

Illustrative, Front Office and Back Office



FRONT OFFICE (Private Banker)

- | | |
|---------------------------------|----------------------------------|
| ✗ Client Meeting Notes | ✓ Needs Discovery |
| ✗ Portfolio Summaries | ✓ Tailored Advisory Provision |
| ✗ Product Information Retrieval | ✓ Client Relationship Management |
| ✗ Internal Documentation | ✓ Commercial Prioritisation |
| ✗ Routine Enquiry Response | ✓ Estate Planning Coordination |

Free up
Resources for



BACK OFFICE (HR Recruiter)

- | | |
|----------------------------|-----------------------------------|
| ✗ CV Screening Support | ✓ Candidate Interview |
| ✗ Candidate Outreach Draft | ✓ Workforce Planning / Advice |
| ✗ Interview Scheduling | ✓ Hiring Judgement |
| ✗ Feedback Summaries | ✓ Talent Event Participation |
| ✗ Candidate Sourcing | ✓ Candidate Relationship Building |

Free up
Resources for



✗ Manual Tasks Led / Augmented by A.I. ✓ Higher-Value Human-led Tasks

WORKFLOW DESIGN (2/2) – EXPLORE NEW WAYS OF WORKING

While adopting a workflow-led approach is essential, ongoing developments in A.I. capabilities also present financial institutions with a unique opportunity to completely reimagine their workflows

A.I. Adoption Pathways

Possible Options




USE CASE PRIORITISATION

A minimum eligibility threshold ensures that institutional resources are focused on viable and compliant use cases, while a robust prioritisation framework determines the optimal delivery sequence to maximise impact

Methodology Overview

Illustrative



Stages

Description

Sample Criteria

Goal

1

SHORTLISTING* (Eligibility Gate)

A pass / fail screening stage that filters out use cases that are not viable for a financial institution to pursue

- Data readiness
- Risk sensitivity threshold

Ensures only feasible and compliant use cases enter the pipeline, **working on what is actually buildable and permissible**

2

PRIORITISATION (Delivery Ordering)

A subsequent comparative assessment stage that ranks shortlisted use cases to determine what to do first, next, and later

- Value impact magnitude
- Value impact reach
- Time-to-market

Ensures that financial institutions are building towards what matters first to **maximise impact and momentum**

← Using a minimum shortlisting threshold ensures organisations focus their time and effort on viable and compliant use cases, while a robust prioritisation framework optimises delivery sequencing to drive maximum impact →

*The concept of eligibility gate is further elaborated on Slide 21 as part of the "Govern" section



SECTION 3

GOVERN


QUINLAN
& ASSOCIATES

A.I. GOVERNANCE FRAMEWORK

As the A.I. landscape evolves and industry practices remain in flux, establishing governance provides both a compliance safeguard and a source of competitive advantage

Governance Considerations

Key Factors



Systemic Playbook (Illustrative)	Use Case Charter
STEP 1: Intake Is the use case sufficiently defined for formal assessment?	- Strategic relevance & business problem - Ownership & stakeholders - Scope & intended users
STEP 2: Eligibility & Prioritisation Should the institution invest further efforts in developing the use case?	- Data availability assessment - Risk and control viability - Feasibility & impact assessment
STEP 3: PoC Approval Should controlled experimentation begin?	- PoC hypothesis and scope - Test design and success criteria - Pre-approved data & risk restrictions
STEP 4: Pilot Approval Should the solution be tested in a live under controlled environment?	- Validated functional solution - Pilot operating model - Pilot control framework
STEP 5: Production Approval Is the solution ready for business-as-usual deployment?	- Risk and control approval - Business-as-usual operating readiness - Production resilience and monitoring
STEP 6: Scale Approval Should the solution be expanded across the organisation?	- Adoption and operating evidence - Realised business benefits - Scalability assessment
STEP 7: Periodic Review Should the use case continue, be modified, restricted, or retired?	- Performance and risk outcomes - Economic and operational impact - Continued relevance and fitness

 Elaborated on the next slide

ELIGIBILITY BEFORE PRIORITISATION

As part of a systematic A.I. playbook, financial institutions should establish non-negotiable eligibility gates for initial shortlisting, only applying feasibility and impact criteria to prioritise use cases that have passed those gates

Shortlisting and Prioritisation Framework

Illustrative


ELIGIBILITY GATE: INITIAL SHORTLISTING			DELIVERY ORDERING: PRIORITISATION OF SHORTLISTED USE CASES	
Criteria	Threshold	Rationale	Weight	Rationale for Weighting
↑ FEASIBILITY ↓	Data Readiness: Is there sufficient data available, and is it structured and of quality to train and run an A.I. model?	>80% ...of the required data is digitised, structured, and owned. Without this, no use case can be built, so the threshold is a hard pass / fail check	15%	Foundational requirement, as poor data quality often prevents A.I. use cases from progressing beyond concept
	Risk Sensitivity: Does the use case touch a process that is subject to regulatory scrutiny and pose material risks?	100% ...of explainability and audit trail requirements can be met, which is a gated requirement for regulated financial institutions	15%	Critical for financial institutions, where customer, conduct, regulatory, model, cyber, and reputational risks are material
	Resource Needs: How much monetary and human resources are needed for the proposed use case?	If a use case does not pass the eligibility gate, it should not proceed to further assessment, ensuring that time and resources are not spent on unsuitable opportunities	10%	Important for feasibility, but should not outweigh strategic relevance, risk, or value creation
	Time-to-Market: Can a working prototype or pilot be delivered quickly and seamlessly?		5%	Early wins build confidence, but speed must remain balanced against governance and control
	Technical Readiness: Could the system (e.g., integration, security, and controls) facilitate this use case?		10%	Existing infrastructure, integration, security, and controls can materially delay or block implementation
↑ IMPACT ↓	Strategic Alignment: Does the initiative directly support the organisation to achieve statutory objectives?		15%	Ensures A.I. initiatives support business strategy, operating priorities, and value creation
	Value Impact Magnitude: How large is the expected benefit (e.g., regulatory quality, faster turnaround)?		15%	Directly measures expected benefits, including revenue, cost, productivity, risk, and customer outcomes
	Value Impact Reach: How widely shared is the use case across functions and divisions?	10%	Scalable use cases can generate broader enterprise-wide benefits across functions, products, and geographies	
	Change Readiness: How open and willing are the broader team to adopt A.I.-assisted ways of working?	5%	Important for adoption, but better addressed through training, change management, and workflow redesign	
TOTAL WEIGHT			100%	

WORKFLOW-LEVEL CONTROL MEASURES

Financial institutions should also extend governance to the workflow level by defining runtime controls, human checkpoints, escalation triggers, resilience mechanisms, and recovery procedures that establish the operating boundaries for implementation

Control Measures

Key Elements, Illustrative Example




HUMAN CHECKPOINT

Defined points in the workflow where human intervention is required, whether it be for approval, escalation, or fallback

Best Practices:

- ✓ Adopt a “design for failure” approach (e.g., log for compliance, escalation triggers, etc.)
- ✓ Define clear trigger conditions for intervention (e.g., risk score, confidence threshold, etc.)
- ✓ Ensure right-sizing of human involvement

An enforceable control mechanism
with **clear accountability**



SECTION 4

ENABLE

SECTION 4-1


TECHNOLOGY

MODEL DEPLOYMENT APPROACH

Once governance establishes permissible operating boundaries, financial institutions must determine whether to use frontier, open-source, or open-weight models, or adopt a multi-model architecture supported by an LLM orchestration layer

Deployment Approach

Frontier vs. Open Source vs. Hybrid



MODEL SELECTION

When selecting an appropriate LLM, a structured scoring methodology across key decision criteria should be employed, adjusting the weighting of each criterion to reflect the needs and sensitivity of different business units

Model Selection Methodology

Key Criteria and Assigned Weights (%)


		Frontier LLM (Closed Source)	Self-hosted LLM (Open Weight/Source)	Front-office (Weight)	Back-office (Weight)
BUSINESS	Data Privacy: How sensitive are the data inputs and conversations processed by the LLM?	■ □ □ □ □	■ ■ ■ ■ ■	20%	5%
	Domain Finetuning: To what extent should the LLM be fine-tuned for the financial institution's business-specific requirements?	■ □ □ □ □	■ ■ ■ ■ ■	20%	5%
	Counterparty Risks: How critical is it for the financial institution to control the model lifecycle and steer clear of vendor lock-in?	■ ■ □ □ □	■ ■ ■ ■ ■	10%	5%
MODEL	Capability: How important is long-context understanding and advanced reasoning for achieving quality outcomes?	■ ■ ■ ■ ■	■ ■ ■ □ □	20%	20%
	Latency: Is there a hard requirement for response time, throughput, uptime, and fallback?	■ ■ ■ ■ ■	■ □ □ □ □	10%	5%
IMPLEMENTATION	Time-to-Market: What is the financial institution's internal timeline for the LLM to be deployed into production environments?	■ ■ ■ ■ ■	■ ■ □ □ □	5%	20%
	Total Cost of Ownership: What is the financial institution's appetite to shoulder the upfront and ongoing cost (e.g., API calls)?	■ ■ ■ □ □	■ ■ ■ □ □	5%	30%
	Operational Readiness: Does the financial institution have internal capabilities (e.g., infrastructure and talents, etc.) to support usage?	■ ■ ■ ■ ■	■ ■ □ □ □	5%	10%
SUITABLE MODEL				Self-hosted LLM	Frontier LLM

FOUNDATIONS OF A.I. ENGINEERING

The success of A.I. deployment depends not only on the model itself, but also on the underlying model-level harness engineering, which must be robust to ensure model capability is reliably translated into production-grade outcomes

A.I. Solution Pillars

Key Components



PROMPT ENGINEERING (2022)

The natural language instruction or question to trigger a specific response

→ *What does the A.I. model need to do?*



CONTEXT ENGINEERING (2024)

The surrounding context and data to ground the A.I.'s response in relevant facts

→ *What does the model need to know?*



HARNESS ENGINEERING (2026)

The infrastructure around a model to automate, control, and validate its end-to-end execution

→ *How should A.I. perform the task?*


A.I. ENGINEERING

Across the different generations of LLM-based solutions being explored, multiple A.I. engineering components must be appropriately configured to ensure that the underlying technology performs reliably within defined governance boundaries

A.I. Solutions Evolution

Applicability of Prompt, Context, and Harness Engineering

✓ Fully Applicable - Applicable to Some Extent ✗ Not Applicable



LEVEL OF AUTONOMY

PROMPT ENGINEERING

To promote consistent and reliable outputs across the organisation, financial institutions should define standardised prompt templates that set out objectives, examples, and step-by-step guidance for employees using GenA.I.-powered applications


Key Components

Well-configured Prompt



SET OBJECTIVES

Designate the model's role and give an overview of the task, as well as relevant contexts to ensure the model captures task-specific nuances



INSTRUCT STEPS

Detail each step involved, including any constraints that the model is prohibited from doing while carrying out the task



INCORPORATE EXAMPLES

Give specific examples of expected outputs help the model learn about the task and the complexity, tone, and concision of expected outputs



USE DELIMITERS (e.g., < >, ' ', "", etc.)

Specify where inputs are inserted in the prompt to help model identify user inputs (e.g., company profiles), especially when they are long



STRUCTURE OUTPUTS

Specify the response with a fixed syntax to parse the output into be more readily available formats to be fed into the following processes

Credit Memo
Client Briefing
AML Case

TASK: Act as a Senior Commercial Credit Analyst. Analyse the raw applicant profile below, perform a credit risk analysis, and generate a standardised credit memo.

ANALYSIS FRAMEWORK
Before formulating your final assessment, you must think step-by-step: (1) Financial Assessment: Analyse liquidity, leverage, and repayment capacity with accounting ratios; (2) Qualitative Analysis: Assess the borrower's industry stability, management experience, and collateral; (3) Risk vs. Mitigant Balancing: Explicitly identify the top 2-3 credit risks and how they are mitigated; (4) Deduction: Formulate the final credit decision based on the balance of these factors.

EXAMPLE
Applicant Profile: Logistics firm operating 5 years. Requesting \$500,000 for fleet expansion; Financials (Annual Revenue: \$4M, Net Profit Margin: 8%. Current Debt Service Coverage Ratio (DSCR): 1.65x (Strong)); Risk Note: Major customer represents 40% of revenue (Concentration Risk); Collateral: Real estate valued at \$700,000.
Analysis Reasoning: (a) Financials: Financial health is strong; cash flow comfortably covers new debt service (1.65x); (b) Qualitative: Experienced management (5 years), but high customer concentration risk makes revenue volatile; (c) Mitigants: The concentration risk is heavily mitigated by strong hard collateral (\$700k real estate) which provides a clear secondary source of repayment.
Output: "executive_summary": "Logistics provider seeking \$500k for fleet expansion. Strong cash flow + solid real estate collateral offset high customer concentration risk.", "key_credit_strengths": ["Strong repayment capacity with a historical DSCR of 1.65x.", "High quality, easily liquidable real estate collateral (\$700k value)."], "key_risks_and_mitigants": ["Risk: Customer concentration (40% of revenue from one client). Mitigant: Long-term contract in place plus substantial asset backing."], "rationale": "The borrower exhibits stable profitability and strong debt service capability. While the loss of their primary client would hurt, the loan is structurally secure due to conservative loan-to-value on the collateral.", "opinion": "Approve".

Credit Memo
Client Briefing
AML Case

TASK: Act as a Senior Commercial Credit Analyst. Analyse the raw applicant profile below, perform a credit risk analysis, and generate a standardised credit memo.

ANALYSIS FRAMEWORK
Before formulating your final assessment, you must think step-by-step: (1) Financial Assessment: Analyse liquidity, leverage, and repayment capacity with accounting ratios; (2) Qualitative Analysis: Assess the borrower's industry stability, management experience, and collateral; (3) Risk vs. Mitigant Balancing: Explicitly identify the top 2-3 credit risks and how they are mitigated; (4) Deduction: Formulate the final credit decision based on the balance of these factors.

EXAMPLE
Applicant Profile: Logistics firm operating 5 years. Requesting \$500,000 for fleet expansion; Financials (Annual Revenue: \$4M, Net Profit Margin: 8%. Current Debt Service Coverage Ratio (DSCR): 1.65x (Strong)); Risk Note: Major customer represents 40% of revenue (Concentration Risk); Collateral: Real estate valued at \$700,000.
Analysis Reasoning: (a) Financials: Financial health is strong; cash flow comfortably covers new debt service (1.65x); (b) Qualitative: Experienced management (5 years), but high customer concentration risk makes revenue volatile; (c) Mitigants: The concentration risk is heavily mitigated by strong hard collateral (\$700k real estate) which provides a clear secondary source of repayment.
Output: "executive_summary": "Logistics provider seeking \$500k for fleet expansion. Strong cash flow + solid real estate collateral offset high customer concentration risk.", "key_credit_strengths": ["Strong repayment capacity with a historical DSCR of 1.65x.", "High quality, easily liquidable real estate collateral (\$700k value)."], "key_risks_and_mitigants": ["Risk: Customer concentration (40% of revenue from one client). Mitigant: Long-term contract in place plus substantial asset backing."], "rationale": "The borrower exhibits stable profitability and strong debt service capability. While the loss of their primary client would hurt, the loan is structurally secure due to conservative loan-to-value on the collateral.", "opinion": "Approve".

Employee Workflow

Select Task

- Choose the task for which predefined templates have been created with all key elements to be loaded in the chatbox

Input Information

- Provide case-specific information for the model to carry out the specific task, such as client profile and financial position

Generate Output


- Request the model to generate outputs following the format and steps as instructed in the template

SKILLS (1/2) – OVERVIEW

Skills are reusable and scalable procedural prompts that can be used by humans on demand or automatically invoked by the relevant A.I. agent based on specific configuration

Lifecycle of Skill

Origination, Storage, and Usage




SKILLS (2/2) – COST & PERFORMANCE IMPACT

On average across the model lineup, skills increase per-trial token consumption, while reducing per-trial cost in practice as a greater share of tokens shifts towards cheaper input-side overhead and improvements in first-pass accuracy with fewer retries

Cost & Performance Improvement

Skills vs. Without Skills, Performance-focused vs. Efficiency-focused Model



¹Price per 1 million tokens; ²Task-macro success rate by agent configuration

Source: SkillsBench, Quinlan & Associates analysis

Key Reasons

Token Economics



OBJECTIVE FUNCTION

The curated skill provided the model with the code template or algebraic structure across steps / calls that it can then attempt to optimise

→ *Focused intelligence = less waste / attempt*



STRUCTURED PROCEDURE

Skills break a complex task into structured steps, increasing tokens from context-reading and acting but fewer total trials are need to achieve outcome

→ *Higher success rate with fewer trial counts*



REASONING / OUTPUT CAPABILITY

While skills introduce additional low-cost input tokens (e.g., workflows, templates), it reduces expensive output, retries, and reasoning loops

→ *Less costly reasoning & correction cycles*

The skills context allows models to reach a solution with significantly less explanatory work

KNOWLEDGE BASE AND MEMORY


The right context engine should be driven by the intended use case, with RAG excelling at retrieving and grounding responses in the document repositories, while LLM Wiki can embed institutional tone and know-how, making it a powerful source of differentiation

Context Framework


RAG vs. LLM Wiki

● User ● Knowledge Base / Source / A.I. Model ● Output → Manual Flow → Automated Flow ✓ High - Moderate ✗ Low

Retrieval-Augmented Generation (“RAG”)



Large Language Model (“LLM”) Wiki



Ease of Setup and Maintenance

✗

Self-managed deployments require chunking logic, an embedding model, a vector database, and ongoing retrieval evaluation and monitoring

Knowledge Depth & Reasoning Quality

-

Reasoning quality matches the underlying model, but synthesis is bounded by retrieval: the model reasons only over the top-ranked chunks returned per query

Knowledge Freshness & Update Latency

-

Source documents can be updated and re-indexed incrementally, with pipelines routinely automated on an event-driven or scheduled basis

Protection from Hallucination Risk

✓

Answers are grounded in retrieved passages and can carry citations back to source, making claims auditable

✓

No complicated back-end infrastructure is required and users can easily drop files into context and edit them directly

✓

The model sees the entire corpus at once with no chunking losses, enabling cross-document synthesis and nuanced reasoning

-

Updates take effect instantly once the markdown file is edited, but past interactions (if not well governed) can silently propagate user errors

-

In-context content provides a degree of grounding, but the model can confidently generate results that falls outside the wiki


¹Conversion to numerical embeddings in order to run similarity search, also known as a vector

AGENTIC LOOP

When the components of harness engineering work together, they enable the agentic loop: the core mechanism by which Agentic A.I. iteratively executes, evaluates, and refines its own outputs autonomously to achieve defined objectives

Agentic Loop

Diagram



Agentic Loop

The autonomous cycle that the model executes to complete a specific task (including observing and matching output with the required conditions to instigate a refinement loop, if needed), all without human intervention

Powered by



Schema

The rulebook that defines the plan of attack and target required output that guides how agents should behave



Orchestration

The coordination of autonomous agents that reason and act with specific control logic and workflow sequencing



Tools

Solutions (e.g., API and MCP) that connect the model to external databases and applications

CONSEQUENCES OF NOT GETTING IT RIGHT

Without effective engineering, A.I. solutions can erode the value of institutional investment, resulting in economic waste, unreliable outputs, and weakened system integrity, resulting in heightened operational, legal, and reputational risks

Key Consequences

Use of A.I. in Other Industries



ECONOMIC WASTE

Runaway Cost from Context Rot

Without disciplined skill and workflow design, A.I. systems reprocess bloated context at every stage and consume far more tokens

Lost Staff Time from Manual Rework

Vague skill definitions and undocumented failure modes renders outputs unusable, forcing employees to re-run the same task or redo it manually



UNRELIABLE OUTPUTS

Hallucination and Factual Errors

Without skills in cross-verifying facts, users may accept A.I.-generated fabrications, leading to compliance, legal, or technical issues

Poor Judgement and Reasoning

Outputs may fail to reflect sound judgment, domain expertise, or organisational intent due to flawed decision logic / workflow



COMPROMISED INTEGRITY

Data Breach

Users frequently leak proprietary information into the A.I. open context if restrictions and data boundaries are not clearly defined in the prompt

External Manipulation

Without strict isolation layers, an A.I. can process malicious instructions (prompt injections) hidden within external sites, compromising systems

J.P.Morgan

Erosion of returns from ungoverned A.I. Usage

J.P. Morgan's Chief Data and Analytics Officer noted that some employees are incurring A.I. token expenses exceeding their salaries, raising skepticism about the true efficacy and return on A.I.

Klarna

Poor judgement in over-automated A.I. support

Klarna's A.I. assistant applied the same automation logic across all query types and hence could not distinguish routine requests from high-stakes cases such as fraud or disputed charges

Community Bank

Sensitive customer data exposed via A.I. tool

A US bank employee uploaded customer data, including Social Security numbers, into an unsanctioned A.I. application, triggering severe regulatory and operational consequences

SECTION 4-2


PEOPLE & PROCESSES

TEAM SETUP

Every role is becoming more managerial in nature, as team members direct and review A.I. outputs, managers orchestrate mixed human and A.I. teams, and leaders steer institutions through a transformation they may not yet fully understand

Team Structure

Traditional Team vs. A.I.-enabled Team



Source: Quinlan & Associates analysis

SENIOR LEADERS

TRUST STEWARD

Senior leadership sets direction and relies on technical teams to recommend and manage technology decisions, given their domain expertise

DIAGNOSTICIAN

Senior leadership is accountable for how A.I. is used across the organisation, the tone on responsible A.I. use, and leading teams who are often more A.I.-capable

MIDDLE MANAGERS

COORDINATOR

Middle management allocates work to team members based on capacity and skills, reviews outputs, and manages performance

HUMAN / A.I. ORCHESTRATOR

Workflow design is not a core mgmt. skill, and middle management owns morale, ensuring team members feel secure, valued, and clear on changes in roles

FRONTLINE STAFF

EXECUTION SPECIALIST

Team members act as skilled contributors, producing output directly and reporting to the team heads, with their performance measured

QUALITY ASSURER

Team members decide what to prompt, review A.I. output critically, and make adaptations, along with adding a judgment layer that may be missing in A.I.

NEW ROLES IN THE A.I. AGE

Senior leaders set strategic direction and investment boundaries, middle managers translate those priorities into workflow-level redesign, and frontline staff determine what works in day-to-day execution

Updated Roles

Key Questions and Purpose

UPDATED ROLES IN THE ORGANISATION	KEY QUESTION TO ADDRESS	PURPOSE
SENIOR LEADERS [DIAGNOSTICIAN] Define the strategic objective, allocate resources, and set risk boundaries, but should not be expected to diagnose every workflow-level inefficiency	• What business outcome are we trying to achieve?	✓ Ensures A.I. starts with a business objective, not a tool
	• What is our level of commitment and preparedness?	✓ Clarifies budget, sponsorship, and risk appetite
	• What does success look like?	✓ Defines measurable KPIs (e.g., revenue uplift, cost reduction, better client experience)
	• What should not be compromised?	✓ Sets non-negotiables (e.g., client trust, compliance, confidentiality, quality control, etc.)
MIDDLE MANAGERS [ORCHESTRATOR] Translate senior-level priorities into workflow redesign by identifying where processes break down, where A.I. should be embedded, and where human control must remain	• Which end-to-end workflow is underperforming?	✓ Identifies friction points (e.g., approval delays) of the specific workflow to redesign
	• Which steps should be automated, augmented, or left human-owned?	✓ Prevents blunt automation of tasks that require judgement
	• Where should human review and approval sit?	✓ Builds control into the redesigned workflow
	• Where do exceptions, workarounds, and edge cases occur?	✓ Ensures the redesigned workflow handles real-world complexity
FRONTLINE STAFF [QUALITY ASSURER] Provide the operational reality by showing where work actually breaks, where A.I. helps or fails, and what must be refined before the workflow can scale	• What actually happens day to day?	✓ Captures operational reality rather than assumed process design
	• Which steps are most repetitive, manual, or frustrating?	✓ Reveals high-potential use cases for A.I. deployment
	• Where is A.I. output useful, unreliable, or unusable?	✓ Calibrates the quality and practicality of the A.I. deployment
	• What needs to be fine-tuned before the workflow can scale?	✓ Creates the feedback loop for adoption and continuous improvement

KEY CONSIDERATIONS FOR SENIORS LEADERS

As A.I. blurs the boundaries between strategy, operations, and technology, senior leaders must understand the underlying shifts, build first-hand capabilities, and actively lead the resulting organisational redesign

Key Considerations

Cheat sheet

	THE SHIFT: WHAT IS AHEAD?	THE IMPACT: WHY DOES THIS MATTER TO YOU?	YOUR MOVE: WHAT SHOULD YOU DO NEXT?
STEP 1 UNDERSTAND	Answering → Acting: A.I. is no longer merely answering your questions, it acts on your behalf Adoption → Scale: Everyone now uses A.I., but only a few convert usage into value through workflow redesign Specify & commission → Sketch & prototype: Business can directly pilot test its own solutions without waiting for IT team	The emergence of new risk category: A system that acts can now do the wrong thing, not just say it The race has started: Having A.I. is no longer a differentiator; advantage now sits in addressing the adoption-to-value gap The new bottleneck: From IT capacity to your own ability to articulate, test, and iterate ideas	Supervise by design: Treat A.I. agents as digital workers: for each one, define who oversees them and how they are audited Fewer deeper bets: Place capital on deployments that redesign whole workflows with clear business case Rewrite the IT contract: Business builds and iterates; technology certifies, validates, and hosts
STEP 2 SPONSOR	Arm's-length sponsorship → First-hand experience: You can no longer sponsor A.I. credibly from a distance Endorsing → Judging: The scarce executive skill becomes judging the output quality for verification and calibration Speed → Fit: Value no longer comes from adopting fast, but from systematically screening where A.I. fits your business	The credibility gap: Sponsors who have never built anything misjudge the potential, the risk, and the speed of deployment Error is a design parameter: Inaccuracy is a commonly reported A.I. output failure, presenting real business risks Focus on high-value work: Judgement-heavy synthesis pay best; deterministic high-volume work often does not	Build one thing: Create one working prototype from your own work in a sandbox before asking others to do the same Verify, then trust: Adopt structured verification habits: demand sources and pressure test cases Rank your P&L: Screen your own function and produce a ranked use-case shortlist before committing any budget
STEP 3 LEAD	Forecasting → Option management: Given the pace of A.I. development, organisational strategy remain open questions Sequential approvals → Risk-tiered gating: Low-risk gates are automated, human review is reserved for high risk Buying software → Buying capability: You are no longer procuring tools, you are buying capability that gets outdated fast	Single scenario does not work: Betting on a single rigid strategy and architecture will lead to irreversible consequences Traditional controls measures are outpaced: Your control framework was built for human-paced change, not for A.I. “Lock-in” is the new risk: Long commitments are dangerous while capability and pricing shift this fast	Stay reversible: Hold options: configurable human gates, reskilling over restructuring, audit-by-design from day one Redesign, no bypass: Classify use cases by risk, automate the low-risk, and concentrate human review on the high-risk Move away from traditional cycle: Transition from annual SaaS subscription to tailored solution leveraging LLM

TRADITIONAL APPROACHES NO LONGER WORK

A.I. transformation cannot rely on a top-down model in which management selects software and pushes it across the organisation, as meaningful value is identified closer to the frontline, where workflows, pain points, and domain knowledge are most visible

Traditional vs Modern Approach

Why it fails and works

TRADITIONAL TOP-DOWN APPROACH

A predefined software solution is selected first, then the organisation attempts to standardise work around the tool

SENIOR LEADERS

- Approve a technology solution based on strategic ambition, vendor claims, and pressure to modernise

MIDDLE MANAGERS


- Convert the selected system into an implementation plan, often with limited ability to reshape the underlying workflow

FRONTLINE STAFF

- Are expected to adopt a fixed workflow, even when the tool does not reflect exceptions, judgement, or informal workarounds

✗ WHY THIS FAILS

Technology defines the redesign boundary, causing the organisation to optimise around system configuration rather than workflow requirements



MODERN BOTTOM-UP APPROACH

A.I. act as flexible workflow discovery and augmentation tools, helping teams expose, test, and codify how work should change

FRONTLINE STAFF

- Use A.I. to surface recurring pain points, draft outputs, test prompts, and reveal where work is repetitive or judgement-heavy

MIDDLE MANAGERS

- Translate frontline insight into A.I.-embedded workflow designs, control points, data requirements, and scalable use cases

SENIOR LEADERS

- Prioritise, fund, govern, and scale the use cases that are strategically relevant, feasible, and measurable

✓ WHY IT WORKS

The workflow defines the technology need, allowing A.I. to be embedded where it improves execution, decision quality, and scalability




IMPORTANCE OF DOMAIN EXPERTISE FROM FRONTLINE STAFF

As A.I. applications progress from simple workflow automation to more advanced A.I.-enabled redesign, transformation quality depends on how effectively human expertise is codified into context, rules, controls, and validation logic

Complementary Relationship

Human Domain Expertise & A.I. Execution Capabilities



Types of Activities	Description	Examples	Goals
1. Documentation (By frontline staff)	Human domain expertise is captured from documents, procedures, examples, business rules, and operational experience	- Business intent - Real-life workflow and example - Judgement criteria	Raw expertise base that explains what matters and why
2. Codification (By A.I.)	A.I. helps structure, clarify, test, and convert domain expertise into repeatable workflow logic	- Codified SOP - Dynamic workflow development - Decision rules & control thresholds	Codified logic that can be consistently applied and governed
3. Automation (By A.I.)	A.I. applies the codified expertise to automate workflow steps within defined permissions and boundaries.	- Retrieval - Drafting - Analysis and routing	Faster and more consistent execution of repeatable manual work
4. Refinement (By frontline staff)	Human experts review outputs, correct errors, handle exceptions, and refine the workflow based on real-world performance	- Output review - Exception handling - Rule refinement	Improved expertise base and stronger next-cycle automation

SECTION 5

HOW WE CAN HELP


QUINLAN
& ASSOCIATES

HOW WE CAN HELP

Before pursuing further A.I. initiatives, particularly agentic A.I., financial institutions should conduct a structured readiness diagnostic to identify gaps, prioritise interventions, and establish a 12–24-month roadmap for scalable adoption

How We Can Help

End-to-End Support



Collaborate with **different functions / divisions** and facilitate **senior stakeholder alignment via workshops** to validate findings, surface priorities, and agree on recommendations to drive A.I. to scale across the organisation



STRATEGY WITH A DIFFERENCE

Copyright © 2026 Quinlan & Associates.

All rights reserved. This report may not be distributed, in whole or in part, without the express written consent of Quinlan & Associates. Quinlan & Associates accepts no liability whatsoever for the actions of third parties in this respect.

The information and opinions in this report were prepared by Quinlan & Associates. This report is not financial or investment advice and should not be relied upon for such advice or as a substitute for professional accounting, tax, legal or financial advice. Quinlan & Associates has made every effort to use reliable, up-to-date and comprehensive information and analysis in this report, but all information is provided without warranty of any kind, express or implied.

Quinlan & Associates disclaims any responsibility to update the information or conclusions in this report. Quinlan & Associates accepts no liability for any loss arising from any action taken or refrained from as a result of information contained in this report or any reports or sources of information referred to herein, or for any consequential, special or similar damages even if advised of the possibility of such damages. This report is not an offer to buy or sell securities or a solicitation of an offer to buy or sell securities.

Website www.quinlanandassociates.com
Email enquiries@quinlanandassociates.com
Telephone (+852) 2618 5000
Address Level 20, One International Finance Centre
1 Harbour View Street, Central
Hong Kong

CONTACT US

Benjamin Quinlan

CEO & Managing Partner

T: +852 2618 5000

E: bquinlan@quinlanandassociates.com

Justin Chung

Associate Partner

T: +852 2618 5000

E: jchung@quinlanandassociates.com

Jeanny Ang

Engagement Manager

T: +852 2618 5000

E: jang@quinlanandassociates.com

Moanna Tang

Senior Consultant

T: +852 2618 5000

E: mtang@quinlanandassociates.com